

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import math
import re
from tqdm import tqdm
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.manifold import TSNE
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, normalized_mutual_info_score, log_loss
from sklearn.linear_model import SGDClassifier, LogisticRegression
from collections import Counter, defaultdict
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold, train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import GaussianNB, MultinomialNB
import math
from sklearn.ensemble import RandomForestClassifier
import six
import sys
sys.modules['sklearn.externals.six'] = six
from mlxtend.classifier import StackingClassifier
import xgboost as xgb
import warnings
warnings.filterwarnings("ignore")
```

```
!gdown --id 1RmX5_q6D7rzoXD7nPUM_s8rKEf1KVMDi #training_text.zip download
!gdown --id 1bSQrw5WmDqgI8hBcr8Pflzatx4xCT0Ex #training_variants.zip download
```

```
/usr/local/lib/python3.9/dist-packages/gdown/cli.py:121: FutureWarning: Option `--id` was deprecated in version 4.3.1 and will be removed in 5.0.
  warnings.warn(
Downloading...
From: https://drive.google.com/uc?id=1RmX5\_q6D7rzoXD7nPUM\_s8rKEf1KVMDi
To: /content/training_text.zip
100% 63.9M/63.9M [00:00<00:00, 179MB/s]
/usr/local/lib/python3.9/dist-packages/gdown/cli.py:121: FutureWarning: Option `--id` was deprecated in version 4.3.1 and will be removed in 5.0.
  warnings.warn(
Downloading...
From: https://drive.google.com/uc?id=1bSQrw5WmDqgI8hBcr8Pflzatx4xCT0Ex
To: /content/training_variants.zip
100% 24.8k/24.8k [00:00<00:00, 36.9MB/s]
```

```
!unzip training_text.zip
!unzip training_variants.zip
```

```
Archive: training_text.zip
  inflating: training_text
Archive: training_variants.zip
  inflating: training_variants
```

```
data_var=pd.read_csv('training_variants')
print("No of data points",data_var.shape[0])
print("No of features",data_var.shape[1])
print("Features",data_var.columns.values)
data_var.head()
```

```
No of data points 3321
No of features 4
Features ['ID' 'Gene' 'Variation' 'Class']
```

	ID	Gene	Variation	Class
0	0	FAM58A	Truncating Mutations	1
1	1	CBL	W802*	2
2	2	CBL	Q249E	2
3	3	CBL	N454D	3
4	4	CBL	L399V	4

```
data_text=pd.read_csv('training_text',sep="\\|\\|",engine="python",names=['ID','Text'],skiprows=1)
print("No of data points",data_text.shape[0])
print("No of features",data_text.shape[1])
print("Features",data_text.columns.values)
data_text.head()
```

```
No of data points 3321
No of features 2
Features ['ID' 'Text']

    ID                                     Text
0    0  Cyclin-dependent kinases (CDKs) regulate a var...
1    1  Abstract Background Non-small cell lung canc...
2    2  Abstract Background Non-small cell lung canc...
3    3  Recent evidence has demonstrated that acquired...

import nltk
nltk.download('stopwords')
Stopwords=set(stopwords.words('english'))

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
```

```
def nlp_preprocessing(text,index,column):
    if(type(text) is not int):
        string=" "
        #replace every char with space
        text=re.sub('[^a-zA-Z0-9\n]',' ',text)
        #replace multispaces wtih single space
        text=re.sub('\s+',' ',text)
        #converting text to lower
        text=text.lower()
        for word in text.split():
            #removing stopword
            if word not in Stopwords:
                string+=word+" "
        data_text[column][index]=string

# text
import time
start=time.time()
for index,row in data_text.iterrows():
    #print(index," ", row)
    if(type(row['Text']) is str):
        nlp_preprocessing(row['Text'],index,'Text')
    else:
        print("No text in id",index)
print("time taken",time.time()-start)

No text in id 1109
No text in id 1277
No text in id 1407
No text in id 1639
No text in id 2755
time taken 36.583274841308594
```

```
data=pd.merge(data_var,data_text,on="ID",how="left")
data.head()
```

	ID	Gene	Variation	Class	Text
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variat...
1	1	CBL	W802*	2	abstract background non small cell lung canc...
2	2	CBL	Q249E	2	abstract background non small cell lung canc...
3	3	CBL	N454D	3	recent evidence demonstrated acquired unipare...
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b linea...

```
data[data.isnull().any(axis=1)]
data.loc[data['Text'].isnull(), 'Text']=data['Gene']+" "+data["Variation"]
data[data['ID']==1109]
```

	ID	Gene	Variation	Class	Text
1109	1109	FANCA	S1088F	1	FANCA S1088F

```
#Splittinh train test cv data
y_true=data['Class'].values
data_gene=data.Gene.str.replace('\s+','_')
data_variation=data.Variation.str.replace('\s+','_')

xtrain,x_test,ytrain,y_test=train_test_split(data,y_true,stratify=y_true,test_size=0.2)
x_train,x_cv,y_train,y_cv=train_test_split(xtrain,ytrain,stratify=ytrain,test_size=0.2)

print("No of data points in train",x_train.shape[0])
print("No of data points in cross validate",x_cv.shape[0])
```

```
print("No of data points in test",x_test.shape[0])
```

```
No of data points in train 2124
```

```
No of data points in cross validate 532
```

```
No of data points in test 665
```

```
#plotting distribution of y i's
```

```
train_class_dist=x_train['Class'].value_counts().sort_index()
```

```
test_class_dist=x_test['Class'].value_counts().sort_index()
```

```
cv_class_dist=x_cv['Class'].value_counts().sort_index()
```

```
def plot_dist_yi(dist,txt,d):
```

```
    txt=str(txt)
```

```
    my_colors='rgbkymc'
```

```
    dist.plot(kind='bar')
```

```
    plt.xlabel("Class")
```

```
    plt.ylabel("Data points per class")
```

```
    plt.title("Distribution of yi's in "+txt+" data")
```

```
    plt.grid()
```

```
    plt.show()
```

```
    sorted_yi=np.argsort(dist.values)
```

```
    for i in sorted_yi:
```


```
        print('Number of data points in class', i+1, ':',dist.values[i], '(', np.round((dist.values[i]/d.shape[0]*100), 3), '%')
```

```
plot_dist_yi(train_class_dist,"Train",x_train)
```

```
plot_dist_yi(test_class_dist,"Test",x_test)
```

```
plot_dist_yi(cv_class_dist,"Cross validation",x_cv)
```

Distribution of yi's in Train data



Number of data points in class 8 : 12 (0.565 %)
 Number of data points in class 9 : 24 (1.13 %)
 Number of data points in class 3 : 57 (2.684 %)
 Number of data points in class 5 : 155 (7.298 %)
 Number of data points in class 6 : 176 (8.286 %)
 Number of data points in class 2 : 289 (13.606 %)
 Number of data points in class 1 : 363 (17.09 %)
 Number of data points in class 4 : 439 (20.669 %)
 Number of data points in class 7 : 609 (28.672 %)

Distribution of yi's in Test data



Random Model



#Plot confusion matrix

```
def plot_confusion_matrix(y_test,y_predicted):
    c=confusion_matrix(y_test,y_predicted)
    b=(c/c.sum(axis=0)) #precision
    a=((c.T)/(c.sum(axis=1))).T
```

```
labels=[1,2,3,4,5,6,7,8,9]
```

```
print("Confusion Matrix")
plt.figure(figsize=(20,7))
sns.heatmap(c,annot=True,cmap='YlGnBu',fmt='.3f',xticklabels=labels,yticklabels=labels)
plt.xlabel("Predicted Class")
plt.ylabel("Original Claass")
plt.show()
```

```
print("Precision Matrix (Column sum=1)")
plt.figure(figsize=(20,7))
sns.heatmap(b,annot=True,cmap='YlGnBu',fmt='.3f',xticklabels=labels,yticklabels=labels)
plt.xlabel("Predicted Class")
plt.ylabel("Original Claass")
plt.show()
```

```
print("Recall Matrix (Row sum=1)")
plt.figure(figsize=(20,7))
sns.heatmap(a,annot=True,cmap='YlGnBu',fmt='.3f',xticklabels=labels,yticklabels=labels)
plt.xlabel("Predicted Class")
plt.ylabel("Original Claass")
plt.show()
```




```
test_data_len=x_test.shape[0]
cv_data_len=x_cv.shape[0]
cv_predicted=np.zeros((cv_data_len,9))
for i in range(cv_data_len):
    rand_probs=np.random.rand(1,9)
    cv_predicted[i]=(rand_probs/(rand_probs.sum(axis=1)))[0]
print("Log loss on Cross validation using Random model",log_loss(y_cv,cv_predicted,eps=1e-15))
```


```
test_predicted=np.zeros((test_data_len,9))
for i in range(test_data_len):
    rand_probs=np.random.rand(1,9)
    test_predicted[i]=(rand_probs/(rand_probs.sum(axis=1)))[0]
print("Log loss on Test using Random model",log_loss(y_test,test_predicted,eps=1e-15))
```

```
y_predicted=np.argmax(test_predicted,axis=1)
plot_confusion_matrix(y_test,y_predicted+1)
```

Log loss on Cross validation using Random model 2.5855291431186194
Log loss on Test using Random model 2.5105824416477884
Confusion Matrix




Univariate Analysis : Gene Feature




```
#code for response coding
def get_gv_fea_dict(alpha,feature,df):
    value_count=x_train[feature].value_counts()
    gv_dict=dict()
    for i,denominator in value_count.items():
        vec=[]
        for k in range(1,10):
            cls_cnt=x_train.loc[(x_train['Class']==k)&(x_train[feature]==i)]
            vec.append((cls_cnt.shape[0]+(alpha*10))/(denominator+(alpha*90)))
        gv_dict[i]=vec
    return gv_dict

def get_gv_feature(alpha,feature,df):
    gv_dict=get_gv_fea_dict(alpha,feature,df)
    value_count=x_train[feature].value_counts()
    gv_fea=[]
    for index,row in df.iterrows():
        if row[feature] in dict(value_count).keys():
            gv_fea.append(gv_dict[row[feature]])
        else:
            gv_fea.append([1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9])
    return gv_fea
```



How many categories are there and how they are distributed



```
def univariate_analysis(train,feature):
    unique_genes=train[feature].value_counts()
    print("No of unique genes: ",unique_genes.shape[0])
    print(unique_genes.head())

    #for distribution
    s=sum(unique_genes.values)
    h=unique_genes.values/s
    plt.plot(h,label="histogram of "+feature)
    plt.xlabel("Index of "+feature)
    plt.ylabel("No of occurence")
    plt.legend()
    plt.grid()
    plt.show()

    c=np.cumsum(h)
    plt.plot(c,label="CDF of "+feature)
    plt.legend()
    plt.grid()
    plt.show()
```



```
univariate_analysis(x_train,'Gene')
```

No of unique genes: 225

BRCA1 168

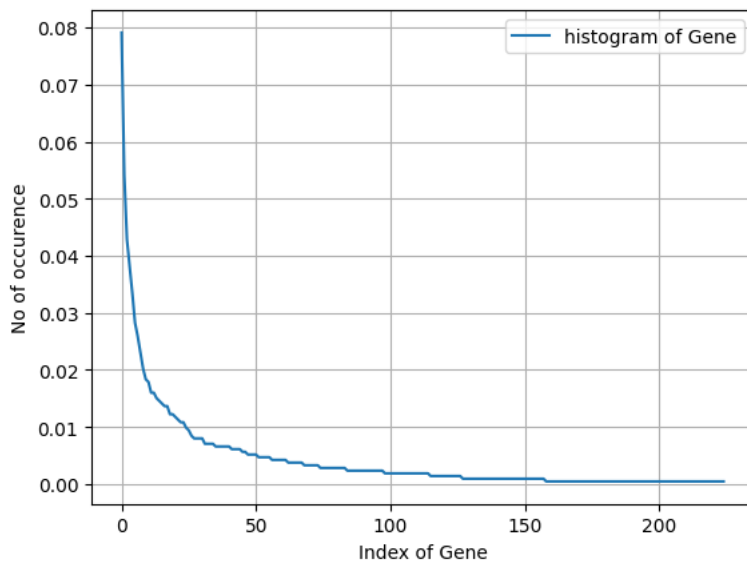
TP53 115

EGFR 91

BRCA2 81

PTEN 71

Name: Gene, dtype: int64



Q] How to featurize gene feature?

Ans One hot encoding and Response coding

alpha=1

train_gene_rc=np.array(get_gv_feature(alpha, 'Gene', x_train))

test_gene_rc=np.array(get_gv_feature(alpha, 'Gene', x_test))

cv_gene_rc=np.array(get_gv_feature(alpha, 'Gene', x_cv))

#ohe

gene_vectorizer=CountVectorizer()

train_gene_ohe=gene_vectorizer.fit_transform(x_train['Gene'])

test_gene_ohe=gene_vectorizer.transform(x_test['Gene'])

cv_gene_ohe=gene_vectorizer.transform(x_cv['Gene'])

type(train_gene_ohe)

scipy.sparse._csr.csr_matrix

Q2] How good is this gene in predicting yi?

Ans Many good ways best is to use a ML model using inly feature yi

```
def predicting_y(alpha, loss, train, cv, test, ytrain, ycv, ytest, weight):
    cv_log_error=[]
    for i in alpha:
        clf=SGDClassifier(alpha=i, penalty='l2', class_weight=weight, loss=loss, random_state=42)
        clf.fit(train, ytrain)
        sig_clf=CalibratedClassifierCV(clf, method='sigmoid')
        sig_clf.fit(train, ytrain)
        predict_y=sig_clf.predict_proba(cv)
        cv_log_error.append(log_loss(ycv, predict_y, labels=clf.classes_, eps=1e-15))
        print("For values of alpha: ", i, " Log loss is: ", log_loss(ycv, predict_y, labels=clf.classes_, eps=1e-15))
```

```
fig, ax=plt.subplots()
ax.plot(alpha, cv_log_error, c='g')
for i, txt in enumerate(np.round(cv_log_error)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
```

```
best_alpha=alpha[np.argmin(cv_log_error)]
clf=SGDClassifier(alpha=best_alpha, penalty='l2', class_weight=weight, loss=loss, random_state=42)
clf.fit(train, ytrain)
sig_clf=CalibratedClassifierCV(clf, method='sigmoid')
sig_clf.fit(train, ytrain)
```

```
predict_y=sig_clf.predict_proba(train)
print("For values of best alpha ", best_alpha, " Train log loss is: ", log_loss(ytrain, predict_y, labels=clf.classes_, eps=1e-15))
predict_y=sig_clf.predict_proba(cv)
```

```

print("For values of best alpha ",best_alpha," Cross Validation log loss is: ",log_loss(ycv,predict_y,labels=clf.classes_,eps=1e-15))
predict_y=sig_clf.predict_proba(test)
print("For values of best alpha ",best_alpha," Test log loss is: ",log_loss(ytest,predict_y,labels=clf.classes_,eps=1e-15))

predicted_y=sig_clf.predict(test)

return predicted_y,best_alpha

```

```


alpha=[10**i for i in range(-5,1)]
predict_y,best_alpha=predicting_y(alpha,'log',train_gene_ohe,cv_gene_ohe,test_gene_ohe,y_train,y_cv,y_test,None)

```

```

For values of alpha: 1e-05 Log loss is: 1.2172652240495068
For values of alpha: 0.0001 Log loss is: 1.1915862037047213
For values of alpha: 0.001 Log loss is: 1.2341761041036268
For values of alpha: 0.01 Log loss is: 1.3365017093525642
For values of alpha: 0.1 Log loss is: 1.421798642293639
For values of alpha: 1 Log loss is: 1.4636580443214573

```



```

For values of best alpha 0.0001 Train log loss is: 1.0027282262203772
For values of best alpha 0.0001 Cross Validation log loss is: 1.1915862037047213
For values of best alpha 0.0001 Test log loss is: 1.1709490349294753

```

Q3] Is Gene feature stable across all data set

Ans Yes it is otherwise cv&test error would be significantly more than train error.

```

print("How many data points are covered by CV and test data are covered by Gene in train data?\nANS:")
test_coverage=x_test[x_test["Gene"].isin(list(set(x_train["Gene"])))].shape[0]
cv_coverage=x_cv[x_cv["Gene"].isin(list(set(x_train["Gene"])))].shape[0]

print("1. In Test data",test_coverage,"out of",x_test.shape[0],":",(test_coverage/len(x_test))*100)
print("2. In CV data",cv_coverage,"out of",x_cv.shape[0],":",(cv_coverage/len(x_cv))*100)

```

How many data points are covered by CV and test data are covered by Gene in train data?

ANS:

1. In Test data 632 out of 665 : 95.03759398496241
2. In CV data 512 out of 532 : 96.2406015037594



▼ Univariate Analysis: Variation feature

```

#understanding categories and distribution
univariate_analysis(x_train,"Variation")

```


No of unique genes: 1945
 Truncating Mutations 53
 Deletion 45
 Amplification 44
 Fusions 17
 G12V 3
 Name: Variation, dtype: int64



#featurizing variation feature which is a categorical feature using response coding and one hot encoding

alpha=1

train_var_rc=np.array(get_gv_feature(alpha,'Variation',x_train))

test_var_rc=np.array(get_gv_feature(alpha,'Variation',x_test))

cv_var_rc=np.array(get_gv_feature(alpha,'Variation',x_cv))

#ohe

var_vectorizer=CountVectorizer()

train_var_ohe=var_vectorizer.fit_transform(x_train['Variation'])

test_var_ohe=var_vectorizer.transform(x_test['Variation'])

cv_var_ohe=var_vectorizer.transform(x_cv['Variation'])


#Understanding how good is variation feature in predicting yi using Ml model

alpha=[10**i for i in range(-5,1)]

predict_y,best_alpha=predicting_y(alpha,'log',train_var_ohe,cv_var_ohe,test_var_ohe,y_train,y_cv,y_test,None)

For values of alpha: 1e-05 Log loss is: 1.6789249057174371
 For values of alpha: 0.0001 Log loss is: 1.6743116975416907
 For values of alpha: 0.001 Log loss is: 1.6834036779603607
 For values of alpha: 0.01 Log loss is: 1.6943987886477914
 For values of alpha: 0.1 Log loss is: 1.7139147611101984
 For values of alpha: 1 Log loss is: 1.7159269275280837

Cross Validation Error for each alpha



For values of best alpha 0.0001 Train log loss is: 0.7213166586877113
 For values of best alpha 0.0001 Cross Validation log loss is: 1.6743116975416907
 For values of best alpha 0.0001 Test log loss is: 1.675487231072414

print("How many data points are covered by CV and test data are covered by Variation in train data?\nANS:")

test_coverage=x_test[x_test["Variation"].isin(list(set(x_train["Variation"])))].shape[0]

cv_coverage=x_cv[x_cv["Variation"].isin(list(set(x_train["Variation"])))].shape[0]

```
print("1. In Test data",test_coverage,"out of",x_test.shape[0],":",(test_coverage/len(x_test))*100)
print("2. In CV data",cv_coverage,"out of",x_cv.shape[0],":",(cv_coverage/len(x_cv))*100)
```

How many data points are covered by CV and test data are covered by Variation in train data?

ANS:

1. In Test data 85 out of 665 : 12.781954887218044
2. In CV data 56 out of 532 : 10.526315789473683

▼ Univariate Analysis : Text feature

```
def extract_dict(cls_text):
    dictionary=defaultdict(int)
    for index,row in cls_text.iterrows():
        for word in row['Text'].split():
            dictionary[word]+=1
    return dictionary

#Building CountVectorizer with all words that occurred min 3 times in train data
#ohe
text_vectorizer=CountVectorizer(min_df=3)
train_text_feature_ohe=text_vectorizer.fit_transform(x_train['Text'])
test_text_feature_ohe=text_vectorizer.transform(x_test['Text'])
cv_text_feature_ohe=text_vectorizer.transform(x_cv['Text'])
```

```
train_text_features=text_vectorizer.get_feature_names_out()
train_text_feat_count=train_text_feature_ohe.sum(axis=0).A1
text_fea_dict=dict(zip(list(train_text_features),train_text_feat_count))
print("Total No of unique words in train data ",len(train_text_features))
```

```
dict_list = []
# dict_list =[] contains 9 dictionaries each corresponds to a class
for i in range(1,10):
    cls_text = x_train[x_train['Class']==i]
    # build a word dict based on the words in that class
    dict_list.append(extract_dict(cls_text))
    # append it to dict_list
```

```
# dict_list[i] is build on i'th class text data
# total_dict is build on whole training text data
total_dict = extract_dict(x_train)
```

```
confuse_array = []
for i in train_text_features:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10)/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)
```

```
#creating response coding for text features
def get_text_rc(df):
    text_feature_rc=np.zeros((df.shape[0],9))
    for i in range(0,9):
        row_index=0
        for index,row in df.iterrows():
            sum_prob=0
            for word in row["Text"].split():
                sum_prob+=math.log(((dict_list[i].get(word,0)+10)/(total_dict.get(word,0)+90)))
            text_feature_rc[row_index][i]=math.exp(sum_prob/len(row["Text"].split()))
            row_index+=1
    return text_feature_rc
```

Total No of unique words in train data 53226

```
train_text_feature_rc=get_text_rc(x_train)
test_text_feature_rc=get_text_rc(x_test)
cv_text_feature_rc=get_text_rc(x_cv)
```

```
#normalize so that row sum is 1
train_text_feature_rc=((train_text_feature_rc.T)/train_text_feature_rc.sum(axis=1)).T
test_text_feature_rc=((test_text_feature_rc.T)/test_text_feature_rc.sum(axis=1)).T
cv_text_feature_rc=((cv_text_feature_rc.T)/cv_text_feature_rc.sum(axis=1)).T
```

```
#normalize ohe
train_text_feature_ohe=normalize(train_text_feature_ohe,axis=0)
test_text_feature_ohe=normalize(test_text_feature_ohe,axis=0)
cv_text_feature_ohe=normalize(cv_text_feature_ohe,axis=0)
```

```
sorted_text_fea_dict=dict(sorted(text_fea_dict.items(),key=lambda x:x[1],reverse=True))
sorted_text_occur=np.array(list(sorted_text_fea_dict.values()))
print(Counter(sorted_text_occur))
```


Counter({3: 4824, 4: 3722, 5: 3202, 6: 2695, 7: 2253, 8: 2023, 9: 1642, 10: 1475, 12: 1237, 11: 1211, 15: 1040, 14: 929, 13: 883, 16: 765, 18: 718

#Understanding how good is text feature in predicting yi using ML model

```
alpha=[10**i for i in range(-5,1)]
```

```
predict_y=predicting_y(alpha, 'log', train_text_feature_ohc, cv_text_feature_ohc, test_text_feature_ohc, y_train, y_cv, y_test, None)
```

```
For values of alpha: 1e-05 Log loss is: 1.3539750727813689
For values of alpha: 0.0001 Log loss is: 1.2097688224393102
For values of alpha: 0.001 Log loss is: 1.1782741206941012
For values of alpha: 0.01 Log loss is: 1.3085488988376641
For values of alpha: 0.1 Log loss is: 1.5015440930080486
For values of alpha: 1 Log loss is: 1.6883298901758286
```



```
For values of best alpha 0.001 Train log loss is: 0.6707982335942583
For values of best alpha 0.001 Cross Validation log loss is: 1.1782741206941012
For values of best alpha 0.001 Test log loss is: 1.1227580272925555
```

```
def get_intersec_text(df):
    df_text_vec=CountVectorizer(min_df=3)
    df_text_fea=df_text_vec.fit_transform(df['Text'])
    df_text_features=df_text_vec.get_feature_names_out()
    df_text_feature_count=df_text_fea.sum(axis=0).A1
    df_text_fea_dict=dict(zip(list(df_text_features),df_text_feature_count))
    len1=len(set(df_text_features))
    len2=len(set(train_text_features) & set(df_text_features))
    return len1,len2

len1,len2=get_intersec_text(x_test)
print(np.round((len2/len1)*100,3),"% of word of test data appeared in train")
len1,len2=get_intersec_text(x_cv)
print(np.round((len2/len1)*100,3),"% of word of cv data appeared in train")

96.846 % of word of test data appeared in train
98.235 % of word of cv data appeared in train
```

▼ ML Model: Data Preparation

```
def predict_plot_confusion_matrix(train_x,train_y,test_x,test_y,clf):
    clf.fit(train_x,train_y)
    sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
    sig_clf.fit(train_x,train_y)
    pred_y=sig_clf.predict(test_x)
    predicted_y=sig_clf.predict_proba(test_x)
    print("log_loss: ",log_loss(test_y,predicted_y))
    print("No of mis classified points:",np.count_nonzero((pred_y-test_y))/test_y.shape[0])
    plot_confusion_matrix(test_y,pred_y)

def report_log_loss(train_x,train_y,test_x,test_y,clf):
    clf.fit(train_x,train_y)
    sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
    sig_clf.fit(train_x,train_y)
    pred_y_prob=sig_clf.predict_proba(test_x)
    return log_loss(test_y,pred_y_prob,eps=1e-15)
```

#Only for Naive Bayes

```
def get_impfeature_names(indices,text,gene,var,no_features):
    gene_count_vec=CountVectorizer()
    var_count_vec=CountVectorizer()
    text_count_vec=CountVectorizer()

    gene_vec=gene_count_vec.fit(x_train["Gene"])
```

```

var_vec=var_count_vec.fit(x_train["Variation"])
text_vec=text_count_vec.fit(x_train["Text"])

fea1_len=len(gene_vec.get_feature_names_out())
fea2_len=len(var_vec.get_feature_names_out())

word_present=0
for i,v in enumerate(indices):
    if(v<fea1_len):
        word=gene_vec.get_feature_names_out()[v]
        yes_no= True if word==gene else False

        if yes_no:
            word_present+=1
            print(i,"Gene feature [",word,"] present in test data point",yes_no)
    elif(v<(fea1_len+fea2_len)):
        word=var_vec.get_feature_names_out()[v-fea1_len]
        yes_no= True if word==var else False
        if yes_no:
            word_present+=1
            print(i,"Variation feature [",word,"] present in test data point",yes_no)

    else:
        word=text_vec.get_feature_names_out()[v-(fea1_len+fea2_len)]
        yes_no= True if word in text.split() else False

        if yes_no:
            word_present+=1
            print(i,"Text feature [",word,"] present in test data point",yes_no)
print("Out of the top ",no_features," features ", word_present, "are present in query point")

```

▼ Stacking Features

```

x_train_ohe=hstack((train_gene_ohe,train_var_ohe,train_text_feature_ohe)).tocsr()
x_test_ohe=hstack((test_gene_ohe,test_var_ohe,test_text_feature_ohe)).tocsr()
x_cv_ohe=hstack((cv_gene_ohe,cv_var_ohe,cv_text_feature_ohe)).tocsr()
print("ONE HOT ENCODING FEATURES:")
print("(Number of datapoints , Number of features) in train data",x_train_ohe.shape)
print("(Number of datapoints , Number of features) in test data",x_test_ohe.shape)
print("(Number of datapoints , Number of features) in cross validation data",x_cv_ohe.shape)

ONE HOT ENCODING FEATURES:
(Number of datapoints , Number of features) in train data (2124, 55422)
(Number of datapoints , Number of features) in test data (665, 55422)
(Number of datapoints , Number of features) in cross validation data (532, 55422)

x_train_rc=np.hstack((train_gene_rc,train_var_rc,train_text_feature_rc))
x_test_rc=np.hstack((test_gene_rc,test_var_rc,test_text_feature_rc))
x_cv_rc=np.hstack((cv_gene_rc,cv_var_rc,cv_text_feature_rc))
print("RESPONSE ENCODING FEATURES:")
print("(Number of datapoints , Number of features) in train data",x_train_rc.shape)
print("(Number of datapoints , Number of features) in test data",x_test_rc.shape)
print("(Number of datapoints , Number of features) in cross validation data",x_cv_rc.shape)

RESPONSE ENCODING FEATURES:
(Number of datapoints , Number of features) in train data (2124, 27)
(Number of datapoints , Number of features) in test data (665, 27)
(Number of datapoints , Number of features) in cross validation data (532, 27)

```

▼ Naive Bayes

```

#Hyperparameter tuning
alpha=[10**i for i in range(-5,4)]
cv_log_error=[]
for i in alpha:
    clf=MultinomialNB(alpha=i)
    clf.fit(x_train_ohe,y_train)
    sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
    sig_clf.fit(x_train_ohe,y_train)
    predict_probs=sig_clf.predict_proba(x_cv_ohe)
    error=log_loss(y_cv,predict_probs,labels=clf.classes_,eps=1e-15)
    cv_log_error.append(error)
    print("For alpha=",i,"Log loss is ",error)

fig,ax=plt.subplots()
ax.plot(np.log10(alpha),cv_log_error,c='g')
for i,txt in enumerate(np.round(cv_log_error,3)):
    ax.annotate((alpha[i],str(txt)),(np.log10(alpha[i]),cv_log_error[i]))
plt.grid()
plt.xticks(np.log10(alpha))
plt.title("Cross validation error for each alpha")
plt.xlabel("Alpha i's")

```

```

plt.ylabel("Error measure")
plt.show()


best_alpha=alpha[np.argmin(cv_log_error)]
#print(best_alpha)
clf=MultinomialNB(alpha=best_alpha)
clf.fit(x_train_ohe,y_train)
sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
sig_clf.fit(x_train_ohe,y_train)

predict_y=sig_clf.predict_proba(x_train_ohe)
print("For best alpha=",best_alpha,"Train Log loss is ",log_loss(y_train,predict_y,labels=clf.classes_,eps=1e-15))
predict_y=sig_clf.predict_proba(x_cv_ohe)
print("For best alpha=",best_alpha,"Cross Validation Log loss is ",log_loss(y_cv,predict_y,labels=clf.classes_,eps=1e-15))
predict_y=sig_clf.predict_proba(x_test_ohe)
print("For best alpha=",best_alpha,"Test Log loss is ",log_loss(y_test,predict_y,labels=clf.classes_,eps=1e-15))


predicted_y=sig_clf.predict(x_test_ohe)
print("No of misclassified points",np.count_nonzero((y_test-predicted_y))/len(y_test))
plot_confusion_matrix(y_test,predicted_y)

```

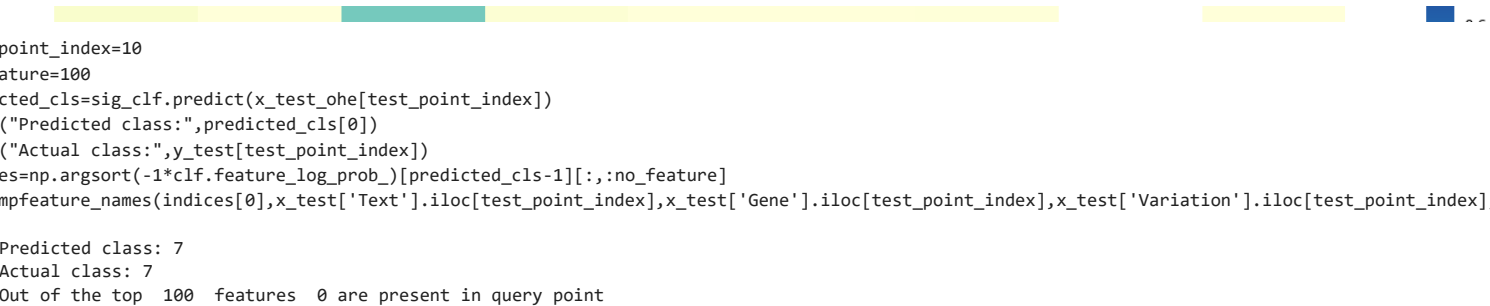
For alpha= 1e-05 Log loss is 1.3304521606331385
For alpha= 0.0001 Log loss is 1.3238240460007502
For alpha= 0.001 Log loss is 1.3149468260580277
For alpha= 0.01 Log loss is 1.3032855263748762
For alpha= 0.1 Log loss is 1.31368828404751
For alpha= 1 Log loss is 1.3401327631194362
For alpha= 10 Log loss is 1.4430863843478463
For alpha= 100 Log loss is 1.3730149031514767
For alpha= 1000 Log loss is 1.3377543228985964



For best alpha= 0.01 Train Log loss is 0.869301546081382
For best alpha= 0.01 Cross Validation Log loss is 1.3032855263748762
For best alpha= 0.01 Test Log loss is 1.2714686384818001
No of misclassified points 0.3924812030075188
Confusion Matrix



Feature importance for correctly classified points



KNN

#in KNN we use response coding than one hot encoding as its high dim, knn doesnt wrk well
alpha=[5,11,15,21,31,41,51,99]

```

cv_log_error=[]
for i in alpha:
    clf=KNeighborsClassifier(n_neighbors=i)
    clf.fit(x_train_rc,y_train)
    sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
    sig_clf.fit(x_train_rc,y_train)
    sig_clf_probs=sig_clf.predict_proba(x_cv_rc)
    cv_log_error.append(log_loss(y_cv,sig_clf_probs,labels=clf.classes_,eps=1e-15))


fig,ax=plt.subplots()
ax.plot(alpha,cv_log_error,c='g')
for i,txt in enumerate(np.round(cv_log_error,3)):
    ax.annotate((alpha[i],str(txt)),(alpha[i],np.round(cv_log_error[i],3)))
plt.grid()
plt.title("CV error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha=alpha[np.argmin(cv_log_error)]
clf=KNeighborsClassifier(n_neighbors=best_alpha)
clf.fit(x_train_rc,y_train)
sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
sig_clf.fit(x_train_rc,y_train)


predict_y=sig_clf.predict_proba(x_train_rc)
print("For best alpha=",best_alpha,"Train Log loss is ",log_loss(y_train,predict_y,labels=clf.classes_,eps=1e-15))
predict_y=sig_clf.predict_proba(x_cv_rc)
print("For best alpha=",best_alpha,"Cross Validation Log loss is ",log_loss(y_cv,predict_y,labels=clf.classes_,eps=1e-15))
predict_y=sig_clf.predict_proba(x_test_rc)
print("For best alpha=",best_alpha,"Test Log loss is ",log_loss(y_test,predict_y,labels=clf.classes_,eps=1e-15))

clf=KNeighborsClassifier(n_neighbors=best_alpha)
predict_plot_confusion_matrix(x_train_rc,y_train,x_test_rc,y_test,clf)


```



For best alpha= 11 Train Log loss is 0.4323979905617689
 For best alpha= 11 Cross Validation Log loss is 1.1454389889449141
 For best alpha= 11 Test Log loss is 1.0984847399634254
 log_loss: 1.0984847399634254
 No of mis classified points: 0.362406015037594
 Confusion Matrix



Precision Matrix (Column sum=1)



Logistic Regression with class balance


```
#hyperparameter tuning
alpha=[10**i for i in range(-6,3)]
predict_y,best_alpha=predicting_y(alpha,'log',x_train_ohc,x_cv_ohc,x_test_ohc,y_train,y_cv,y_test,'balanced')
plot_confusion_matrix(y_test,predict_y)
```


Line graph titled "Cross Validation Error for each alpha". The x-axis is labeled "Alpha i's" and ranges from 0 to 100. The y-axis is labeled "Error measure" and ranges from 1.1 to 1.7. The graph shows a green line representing the cross-validation error for different values of alpha. The error starts at a low value for small alpha and increases sharply as alpha increases, eventually plateauing at a value of 2.0 for alpha values of 10 and above.

Alpha i's	Error measure
0.0001	1.1
0.001	1.12
0.0001	1.15
0.01	1.2
1e-05	1.3
1e-06	1.35
0.1	1.45
1	1.7
10	1.75
100	1.75

[illegible][illegible]

© 2013 Pearson Education, Inc. or its affiliate(s). All rights reserved. Pearson Education, Inc., publishing as Pearson Benjamin Cummings, 101 Philip Drive, Assinippi Park, New York, NY 10984-2135



Feature Importance

```
#Correctly classified points
clf=SGDClassifier(alpha=best_alpha,loss='log',class_weight='balanced',random_state=42,penalty='l2')
clf.fit(x_train_ohc,y_train)
sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
sig_clf.fit(x_train_ohc,y_train)


test_point_index=2
no_feature=500
predicted_cls=sig_clf.predict(x_test_ohc[test_point_index])
print("Predicted class:",predicted_cls[0])
print("Actual class:",y_test[test_point_index])
indices=np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,no_feature]
get_impfeature_names(indices[0],x_test['Text'].iloc[test_point_index],x_test['Gene'].iloc[test_point_index].lower(),x_test['Variation'].iloc[test_point_index])

Predicted class: 1
Actual class: 5
269 Text feature [ a1752p ] present in test data point True
341 Text feature [ data ] present in test data point True
368 Text feature [ 492 ] present in test data point True
441 Text feature [ basal ] present in test data point True
493 Text feature [ developing ] present in test data point True
Out of the top 500 features 5 are present in query point
```

Logistic Regression without class balance

```
#hyperparameter tuning
alpha=[10**i for i in range(-6,3)]
predict_y,best_alpha=predicting_y(alpha,'log',x_train_ohc,x_cv_ohc,x_test_ohc,y_train,y_cv,y_test,None)
plot_confusion_matrix(y_test,predict_y)
```





▼ Feature importance

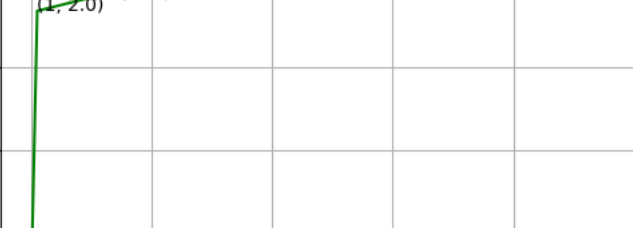
```
#Correctly classified points
clf=SGDClassifier(alpha=best_alpha,loss='log',class_weight=None,random_state=42,penalty='l2')
clf.fit(x_train_ohc,y_train)
sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
sig_clf.fit(x_train_ohc,y_train)

test_point_index=5
no_feature=100
predicted_cls=sig_clf.predict(x_test_ohc[test_point_index])
print("Predicted class:",predicted_cls[0])
print("Actual class:",y_test[test_point_index])
indices=np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,:no_feature]
get_impfeature_names(indices[0],x_test['Text'].iloc[test_point_index],x_test['Gene'].iloc[test_point_index].lower(),x_test['Variation'].iloc[test_point_index])

Predicted class: 7
Actual class: 3
9 Gene feature [ pdgfra ] present in test data point True
Out of the top 100 features 1 are present in query point
```

▼ Linear SVM


```
#hyperparameter tuning
alpha=[10**i for i in range(-5,3)]
predict_y,best_alpha=predicting_y(alpha,'hinge',x_train_ohc,x_cv_ohc,x_test_ohc,y_train,y_cv,y_test,None)
plot_confusion_matrix(y_test,predict_y)
```




A line plot titled "Cross Validation Error for each alpha". The x-axis is labeled "Alpha i's" and ranges from 0 to 100. The y-axis is labeled "Error measure" and ranges from 1.2 to 1.7. A green line shows the error measure for different values of alpha. The line starts at (0, 1.15), rises sharply to (1, 1.67), and then levels off, reaching a maximum error of 1.69 at alpha = 100. Several points are labeled on the line: (0.0001, 1.0), (0.001, 1.0), (0.01, 1.0), (0.1, 1.0), (1e-05, 1.0), (1, 2.0), (10, 2.0), and (100, 2.0).

[illegible][illegible]

1	0.684	0.018	0.000	0.219	0.053	0.018	0.009	0.000	0.000
---	-------	-------	-------	-------	-------	-------	-------	-------	-------



Feature importance



```
#Correctly classified points
clf=SGDClassifier(alpha=best_alpha,loss='hinge',random_state=42,penalty='l2')
clf.fit(x_train_ohc,y_train)
sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
sig_clf.fit(x_train_ohc,y_train)

test_point_index=2
no_feature=100
predicted_cls=sig_clf.predict(x_test_ohc[test_point_index])
print("Predicted class:",predicted_cls[0])
print("Actual class:",y_test[test_point_index])
indices=np.argsort(-1*abs(clf.coef_))[predicted_cls-1][:,no_feature]
get_impfeature_names(indices[0],x_test['Text'].iloc[test_point_index],x_test['Gene'].iloc[test_point_index].lower(),x_test['Variation'].iloc[test_point_index])

Predicted class: 4
Actual class: 5
1 Gene feature [ brca1 ] present in test data point True
Out of the top 100 features 1 are present in query point
```

Random Forest

```
def random_forest(alpha,max_depth,train,cv,test,ytrain,ycv,ytest):
    cv_log_error=[]
    for i in alpha:
        for j in max_depth:
            clf=RandomForestClassifier(n_estimators=i,criterion='gini',max_depth=j,n_jobs=-1)
            clf.fit(train,ytrain)
            sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
            sig_clf.fit(train,ytrain)
            sig_clf_probs=sig_clf.predict_proba(cv)
            error=log_loss(ycv,sig_clf_probs,labels=clf.classes_,eps=1e-15)
            cv_log_error.append(error)
            print("For n-estimators:",i,"and max depth:",j,"Log loss is",error)

    best_alpha = np.argmin(cv_log_error)
    best_n_estimator=alpha[int(best_alpha/len(max_depth))]
    best_max_depth=max_depth[best_alpha%len(max_depth)]
    print("\nBest max depth is:",best_max_depth,"\nBest alpha is:",best_n_estimator)

    clf=RandomForestClassifier(n_estimators=best_n_estimator,max_depth=best_max_depth,criterion='gini',n_jobs=-1)
    clf.fit(train,ytrain)
    sig_clf=CalibratedClassifierCV(clf,method='sigmoid')
    sig_clf.fit(train,ytrain)
    predict_y=sig_clf.predict_proba(train)
    print("\nTrain log loss is:",log_loss(ytrain,predict_y,labels=clf.classes_,eps=1e-15))
    predict_y=sig_clf.predict_proba(cv)
    print("Cross validation log loss is:",log_loss(ycv,predict_y,labels=clf.classes_,eps=1e-15))
    predict_y=sig_clf.predict_proba(test)
    print("Test log loss is:",log_loss(ytest,predict_y,labels=clf.classes_,eps=1e-15),"\n")

    predicted_y=sig_clf.predict(test)
    plot_confusion_matrix(ytest,predicted_y)

    #Feature importance
    test_point_index=2
    no_feature=500
    indices=np.argsort(-clf.feature_importances_)
    print("***10,FEATURE IMPORTANCE",***10)
    get_impfeature_names(indices[:no_feature],x_test['Text'].iloc[test_point_index],x_test['Gene'].iloc[test_point_index].lower(),x_test['Variation'].iloc[test_point_index])


alpha=[100,500,1000,2000]
max_depth=[1,5,10]
random_forest(alpha,max_depth,x_train_ohc,x_cv_ohc,x_test_ohc,y_train,y_cv,y_test)
```

For n-estimators: 100 and max depth: 1 Log loss is 1.4625915842775723
For n-estimators: 100 and max depth: 5 Log loss is 1.2396319203643955
For n-estimators: 100 and max depth: 10 Log loss is 1.1904283937875362
For n-estimators: 500 and max depth: 1 Log loss is 1.4516162863029844
For n-estimators: 500 and max depth: 5 Log loss is 1.211258543143051
For n-estimators: 500 and max depth: 10 Log loss is 1.1671898189556358
For n-estimators: 1000 and max depth: 1 Log loss is 1.4487853784328595
For n-estimators: 1000 and max depth: 5 Log loss is 1.2102655835317433
For n-estimators: 1000 and max depth: 10 Log loss is 1.1657288473474445
For n-estimators: 2000 and max depth: 1 Log loss is 1.4534803311160043
For n-estimators: 2000 and max depth: 5 Log loss is 1.2124605448663515
For n-estimators: 2000 and max depth: 10 Log loss is 1.1621057756141928


Best max depth is: 10
Best alpha is: 2000

Train log loss is: 0.6671612651248321
Cross validation log loss is: 1.1618550416329043
Test log loss is: 1.1640788252181422


Confusion Matrix




Precision Matrix (Column sum=1)



Recall Matrix (Row sum=1)





***** FEATURE IMPORTANCE *****
18 Text feature [effects] present in test data point True
27 Text feature [c1697r] present in test data point True
36 Text feature [allowed] present in test data point True
140 Gene feature [brca1] present in test data point True
203 Text feature [accordance] present in test data point True
384 Text feature [350] present in test data point True
Out of the top 500 features 6 are present in query point


```
#Random forest with response coding
alpha=[100,500,1000,2000]
max_depth=[3,5,10]
random_forest(alpha,max_depth,x_train_rc,x_cv_rc,x_test_rc,y_train,y_cv,y_test)
```


For n-estimators: 100 and max depth: 3 Log loss is 1.529411766106504
For n-estimators: 100 and max depth: 5 Log loss is 1.3837640968298808
For n-estimators: 100 and max depth: 10 Log loss is 1.6354660232739784
For n-estimators: 500 and max depth: 3 Log loss is 1.5443849070401674
For n-estimators: 500 and max depth: 5 Log loss is 1.382440791404571
For n-estimators: 500 and max depth: 10 Log loss is 1.6797595421433067
For n-estimators: 1000 and max depth: 3 Log loss is 1.554031928052659
For n-estimators: 1000 and max depth: 5 Log loss is 1.3803494134925116
For n-estimators: 1000 and max depth: 10 Log loss is 1.6746746421490202
For n-estimators: 2000 and max depth: 3 Log loss is 1.5240792423713736
For n-estimators: 2000 and max depth: 5 Log loss is 1.3976678133839793
For n-estimators: 2000 and max depth: 10 Log loss is 1.690998607389119


Best max depth is: 5
Best alpha is: 1000

Train log loss is: 0.05828111040156028
Cross validation log loss is: 1.398814190202243
Test log loss is: 1.3523510797104281

Confusion Matrix



Precision Matrix (Column sum=1)



Stack the Models

```
Predicted Class

clf1=SGDClassifier(alpha=0.001,penalty='l2',loss='log',class_weight='balanced',random_state=42)
clf1.fit(x_train_oh,y_train)
sig_clf1=CalibratedClassifierCV(clf1,method='sigmoid')
sig_clf1.fit(x_train_oh,y_train)
clf2=SGDClassifier(alpha=0.01,penalty='l2',loss='hinge',class_weight='balanced',random_state=42)
clf2.fit(x_train_oh,y_train)
sig_clf2=CalibratedClassifierCV(clf2,method='sigmoid')
sig_clf2.fit(x_train_oh,y_train)
clf3=MultinomialNB(alpha=0.01)
clf3.fit(x_train_oh,y_train)
sig_clf3=CalibratedClassifierCV(clf3,method='sigmoid')
sig_clf3.fit(x_train_oh,y_train)

print("Logistic Regression log loss is",log_loss(y_cv,sig_clf1.predict_proba(x_cv_oh)))
print("Linear SVM log loss is",log_loss(y_cv,sig_clf2.predict_proba(x_cv_oh)))
print("Naive Baye's log loss is",log_loss(y_cv,sig_clf3.predict_proba(x_cv_oh)))

alpha=[0.0001,0.001,0.01,0.1,1,10]
```


```

best_alpha=999
for i in alpha:
    lr=LogisticRegression(C=i)
    sclf=StackingClassifier(classifiers=[sig_clf1,sig_clf2,sig_clf3],meta_classifier=lr,use_probab=True)
    sclf.fit(x_train_oh,y_train)
    print("Stacking Classifier log loss for alpha=",i,"is",log_loss(y_cv,sclf.predict_proba(x_cv_oh)))
    log_error=log_loss(y_cv,sclf.predict_proba(x_cv_oh))
    if best_alpha>log_error:
        best_alpha=log_error


#using best_alpha
lr=LogisticRegression(C=best_alpha)
sclf=StackingClassifier(classifiers=[sig_clf1,sig_clf2,sig_clf3],meta_classifier=lr,use_probab=True)
sclf.fit(x_train_oh,y_train)
log_error=log_loss(y_train,sclf.predict_proba(x_train_oh))
print("Train log loss is",log_error)
log_error=log_loss(y_cv,sclf.predict_proba(x_cv_oh))
print("Cross Validation log loss is",log_error)
log_error=log_loss(y_test,sclf.predict_proba(x_test_oh))
print("Test log loss is",log_error)
plot_confusion_matrix(y_test,sclf.predict(x_test_oh))

```

Logistic Regression log loss is 1.1114907197534853
Linear SVM log loss is 1.1482773887354718
Naive Baye's log loss is 1.3032855263748762
Stacking Classifier log loss for alpha= 0.0001 is 1.813936459418986
Stacking Classifier log loss for alpha= 0.001 is 1.6895099605075667
Stacking Classifier log loss for alpha= 0.01 is 1.2753604576490178
Stacking Classifier log loss for alpha= 0.1 is 1.1962510192121893
Stacking Classifier log loss for alpha= 1 is 1.4486500113475416
Stacking Classifier log loss for alpha= 10 is 1.825199899541283
Train log loss is 0.28578335316326897
Cross Validation log loss is 1.4703482029288244
Test log loss is 1.2944589119385688
Confusion Matrix



Precision Matrix (Column sum=1)




Maximum Voting Classifier




```
from sklearn.ensemble import VotingClassifier
vcclf=VotingClassifier(estimators=[('lr',sig_clf1),('svc',sig_clf2),('NB',sig_clf3)],voting='soft')
vcclf.fit(x_train_ohc,y_train)
log_error=log_loss(y_train,vcclf.predict_proba(x_train_ohc))
print("Train log loss is",log_error)
log_error=log_loss(y_cv,vcclf.predict_proba(x_cv_ohc))
print("Cross Validation log loss is",log_error)
log_error=log_loss(y_test,vcclf.predict_proba(x_test_ohc))
print("Test log loss is",log_error)

print("No of misclassified points",np.count_nonzero((y_test-vcclf.predict(x_test_ohc)))/y_test.shape[0])
plot_confusion_matrix(y_test,vcclf.predict(x_test_ohc))
```


Train log loss is 0.68294794994484
Cross Validation log loss is 1.0944732978578402
Test log loss is 1.065379018564127
No of misclassified points 0.362406015037594
Confusion Matrix



Precision Matrix (Column sum=1)



Recall Matrix (Row sum=1)



Some Alternatives to reduce Log loss

```
gene=CountVectorizer(ngram_range=(1,2))
x_gene_train=gene.fit_transform(x_train['Gene'])
x_gene_test=gene.transform(x_test['Gene'])
x_gene_cv=gene.transform(x_cv['Gene'])

var=CountVectorizer(ngram_range=(1,2))
x_var_train=var.fit_transform(x_train['Variation'])
x_var_test=var.transform(x_test['Variation'])
x_var_cv=var.transform(x_cv['Variation'])

txt=CountVectorizer(ngram_range=(1,2))
x_txt_train=txt.fit_transform(x_train['Text'])
x_txt_test=txt.transform(x_test['Text'])
```

```
x_txt_cv=txt.transform(x_cv['Text'])
```

```
x_train_oh=hstack((x_gene_train,x_var_train,x_txt_train)).tocsr()
```

```
x_test_oh=hstack((x_gene_test,x_var_test,x_txt_test)).tocsr()
```

```
x_cv_oh=hstack((x_gene_cv,x_var_cv,x_txt_cv)).tocsr()
```

```
print("ONE HOT ENCODING USING N-GRAMS")
```

```
print("Train data",x_train_oh.shape)
```

```
print("Test data",x_test_oh.shape)
```

```
print("CV data",x_cv_oh.shape)
```

```
ONE HOT ENCODING USING N-GRAMS
```

```
Train data (2124, 2347104)
```

```
Test data (665, 2347104)
```


```
CV data (532, 2347104)
```

```
alpha=[10**i for i in range(-5,3)]
```



```
predict_y,best_alpha=predicting_y(alpha,'log',x_train_oh,x_cv_oh,x_test_oh,y_train,y_cv,y_test,'balanced')
```

```
plot_confusion_matrix(y_test,predict_y)
```

For values of alpha: 1e-05 Log loss is: 1.8308895984571567
For values of alpha: 0.0001 Log loss is: 1.8308895984571567
For values of alpha: 0.001 Log loss is: 1.3326321375867771
For values of alpha: 0.01 Log loss is: 1.288324512830357
For values of alpha: 0.1 Log loss is: 1.28071354388059
For values of alpha: 1 Log loss is: 1.1757203605372277
For values of alpha: 10 Log loss is: 1.1559412976670134
For values of alpha: 100 Log loss is: 1.28523377654199



For values of best alpha 10 Train log loss is: 0.8488637344425812
For values of best alpha 10 Cross Validation log loss is: 1.1559412976670134
For values of best alpha 10 Test log loss is: 1.1458739599135879



```
gene=TfidfVectorizer(ngram_range=(1,2))
x_gene_train=gene.fit_transform(x_train['Gene'])
x_gene_test=gene.transform(x_test['Gene'])
x_gene_cv=gene.transform(x_cv['Gene'])
```

```
var=TfidfVectorizer(ngram_range=(1,2))
x_var_train=var.fit_transform(x_train['Variation'])
x_var_test=var.transform(x_test['Variation'])
x_var_cv=var.transform(x_cv['Variation'])
```

```
txt=TfidfVectorizer(ngram_range=(1,2))
x_txt_train=txt.fit_transform(x_train['Text'])
x_txt_test=txt.transform(x_test['Text'])
x_txt_cv=txt.transform(x_cv['Text'])
```

```
x_train_oh=.hstack((x_gene_train,x_var_train,x_txt_train)).tocsr()
x_test_oh=.hstack((x_gene_test,x_var_test,x_txt_test)).tocsr()
x_cv_oh=.hstack((x_gene_cv,x_var_cv,x_txt_cv)).tocsr()
```


```
print("ONE HOT ENCODING USING N-GRAMS")
print("Train data",x_train_oh.shape)
print("Test data",x_test_oh.shape)
print("CV data",x_cv_oh.shape)
```

ONE HOT ENCODING USING N-GRAMS
Train data (2124, 2391985)
Test data (665, 2391985)
CV data (532, 2391985)




```
alpha=[10**i for i in range(-5,3,2)]  
predict_y,best_alpha=predicting_y(alpha,'log',x_train_ohe,x_cv_ohe,x_test_ohe,y_train,y_cv,y_test,'balanced')  
plot_confusion_matrix(y_test,predict_y)
```

For values of alpha: 1e-05 Log loss is: 0.9875390437012223
For values of alpha: 0.001 Log loss is: 1.0183019510333382
For values of alpha: 0.1 Log loss is: 1.3654300888324855
For values of alpha: 10 Log loss is: 1.4698262479224597



For values of best alpha 1e-05 Train log loss is: 0.3190780447067044
For values of best alpha 1e-05 Cross Validation log loss is: 0.9875390437012223
For values of best alpha 1e-05 Test log loss is: 1.004237153037683
Confusion Matrix



```
gene=TfidfVectorizer(ngram_range=(1,3))
x_gene_train=gene.fit_transform(x_train['Gene'])
x_gene_test=gene.transform(x_test['Gene'])
x_gene_cv=gene.transform(x_cv['Gene'])
```

```
var=TfidfVectorizer(ngram_range=(1,3))
x_var_train=var.fit_transform(x_train['Variation'])
x_var_test=var.transform(x_test['Variation'])
x_var_cv=var.transform(x_cv['Variation'])
```

```
txt=TfidfVectorizer(ngram_range=(1,3))
x_txt_train=txt.fit_transform(x_train['Text'])
x_txt_test=txt.transform(x_test['Text'])
x_txt_cv=txt.transform(x_cv['Text'])
```


```
x_train_oh=.hstack((x_gene_train,x_var_train,x_txt_train)).tocsr()
x_test_oh=.hstack((x_gene_test,x_var_test,x_txt_test)).tocsr()
x_cv_oh=.hstack((x_gene_cv,x_var_cv,x_txt_cv)).tocsr()
```

```
print("ONE HOT ENCODING USING TRI-GRAMS")
print("Train data",x_train_oh.shape)
print("Test data",x_test_oh.shape)
print("CV data",x_cv_oh.shape)
```

```
ONE HOT ENCODING USING TRI-GRAMS
Train data (2124, 7063833)
Test data (665, 7063833)
CV data (532, 7063833)
```


```
alpha=[10**i for i in range(-5,3,2)]
predict_y,best_alpha=predicting_y(alpha,'log',x_train_oh,x_cv_oh,x_test_oh,y_train,y_cv,y_test,'balanced')
plot_confusion_matrix(y_test,predict_y)
```


For values of alpha: 1e-05 Log loss is: 0.9882629229794864
For values of alpha: 0.001 Log loss is: 1.01572798009445
For values of alpha: 0.1 Log loss is: 1.3505982950323543
For values of alpha: 10 Log loss is: 1.4496786466163119




For values of best alpha 1e-05 Train log loss is: 0.31623719318539645
For values of best alpha 1e-05 Cross Validation log loss is: 0.9882629229794864
For values of best alpha 1e-05 Test log loss is: 1.0016995798303003


Confusion Matrix




Precision Matrix (Column sum=1)



Recall Matrix (Row sum=1)





```

gene=TfidfVectorizer(ngram_range=(1,2),max_features=100000)
x_gene_train=gene.fit_transform(x_train[ 'Gene' ])
x_gene_test=gene.transform(x_test[ 'Gene' ])
x_gene_cv=gene.transform(x_cv[ 'Gene' ])

var=TfidfVectorizer(ngram_range=(1,2),max_features=100000)
x_var_train=var.fit_transform(x_train[ 'Variation' ])
x_var_test=var.transform(x_test[ 'Variation' ])
x_var_cv=var.transform(x_cv[ 'Variation' ])

txt=TfidfVectorizer(ngram_range=(1,2),max_features=100000)
x_txt_train=txt.fit_transform(x_train[ 'Text' ])
x_txt_test=txt.transform(x_test[ 'Text' ])
x_txt_cv=txt.transform(x_cv[ 'Text' ])

x_train_ohe=hstack((x_gene_train,x_var_train,x_txt_train)).tocsr()
x_test_ohe=hstack((x_gene_test,x_var_test,x_txt_test)).tocsr()
x_cv_ohe=hstack((x_gene_cv,x_var_cv,x_txt_cv)).tocsr()


print("ONE HOT ENCODING USING Max features")
print("Train data",x_train_ohe.shape)
print("Test data",x_test_ohe.shape)
print("CV data",x_cv_ohe.shape)

ONE HOT ENCODING USING Max features
Train data (2124, 102365)
Test data (665, 102365)
CV data (532, 102365)

alpha=[10**i for i in range(-6,3)]
predict_y,best_alpha=predicting_y(alpha,'log',x_train_ohe,x_cv_ohe,x_test_ohe,y_train,y_cv,y_test,'balanced')
plot_confusion_matrix(y_test,predict_y)

```

For values of alpha: 1e-06 Log loss is: 1.1071144489712217
For values of alpha: 1e-05 Log loss is: 0.9804415521085811
For values of alpha: 0.0001 Log loss is: 0.9599657097821842
For values of alpha: 0.001 Log loss is: 1.0176285750187575
For values of alpha: 0.01 Log loss is: 1.1631332698095007
For values of alpha: 0.1 Log loss is: 1.3703570095164703
For values of alpha: 1 Log loss is: 1.462605099614473
For values of alpha: 10 Log loss is: 1.4758190773816635
For values of alpha: 100 Log loss is: 1.4774423842954076



For values of best alpha 0.0001 Train log loss is: 0.3559905571167076
For values of best alpha 0.0001 Cross Validation log loss is: 0.9599657097821842
For values of best alpha 0.0001 Test log loss is: 0.9764762272644417

